

实验三报告

关卡一：openGauss 数据库的编译和安装

1. 关卡验证

步骤 1 首先需要对数据库状态进行验证。

```
[omm@opengauss01 openGauss-server]$ gs_ctl status
```

(截图语句和执行结果)

```
omm@opengauss01 ~$ gs_ctl status
[2022-12-06 15:04:55.356][228668][][gs_ctl]: gs_ctl status,datadir is /opt/software/openGauss/data
gs_ctl: server is running (PID: 228612)
/opt/software/openGauss/bin/gaussdb "-D" "/opt/software/openGauss/data"
omm@opengauss01 ~$
```

步骤 2 对数据库进程进行截图验证，需包含数据库服务器的主机名。

```
[omm@opengauss01 openGauss-server]$ ps -ef|grep omm
```

(截图语句和执行结果)

```
omm@opengauss01 ~$ ps -ef|grep omm
root      228360      5909    0 15:02 pts/0    00:00:00 su - omm
omm       228361    228360    0 15:02 pts/0    00:00:00 -bash
omm       228406    228361    0 15:02 pts/0    00:00:00 zsh
omm       228612          1  1 15:04 pts/0    00:00:01 /opt/software/openGauss/bin/gaussdb -D /opt/software/openGauss/data
omm       228680    228406    0 15:05 pts/0    00:00:00 ps -ef
omm       228681    228406    0 15:05 pts/0    00:00:00 grep --color=auto --exclude-dir=.bzr --exclude-dir=CVS --exclude-dir=.git --exclude-dir=.hg --exclude-dir=.svn --exclude-dir=.idea --exclude-dir=.tox omm
omm@opengauss01 ~$
```

关卡二：openGauss 数据导入及基本操作

1. 关卡验证

步骤 12 登录数据库验证

```
[omm@opengauss01 dbgen]$ gsql -d tpch -p 5432 -r
tpch=# select count(*) from supplier;
```

(截图语句和执行结果)

```
omm@opengauss01 /opt/software/tpch-kit/dbgen git:(master) × sh ./LoadData.sh
Loading customer...
Loading lineitem...
Loading nation...
Loading orders...
Loading partsupp...
Loading part...
Loading region...
Loading supplier...
omm@opengauss01 /opt/software/tpch-kit/dbgen git:(master) × gsql -d tpch -p 5432 -r
gsql ((GaussDB Kernel V500R002C00 build b2ff10be) compiled at 2022-12-06 14:40:54 commit 0 last mr debug)
Non-SSL connection (SSL connection is recommended when requiring high-security)
Type "help" for help.

tpch=# select count(*) from supplier;
count
-----
10000
(1 row)

tpch=#
```

步骤 21 登录数据库进行验证

```
[omm@opengauss01 ~]$ gsql -d tpch -p 5432 -r
tpch=# \dt
```

(截图语句和执行结果)

```
omm@opengauss01 /opt/software/tpch-kit/dbgen/queries git:(master) × gsql -d tpch -p 5432 -r
gsql ((GaussDB Kernel V500R002C00 build b2ff10be) compiled at 2022-12-06 14:40:54 commit 0 last mr debug)
Non-SSL connection (SSL connection is recommended when requiring high-security)
Type "help" for help.

tpch=# \dt

List of relations
Schema | Name | Type | Owner | Storage
-----+-----+-----+-----+-----
public | address_dimension | table | omm | {orientation=row,compression=no}
public | customer | table | omm | {orientation=row,compression=no}
public | date_dimension | table | omm | {orientation=row,compression=no}
public | lineitem | table | omm | {orientation=row,compression=no}
public | litemail_orders | table | omm | {orientation=row,compression=no}
public | nation | table | omm | {orientation=row,compression=no}
public | orders | table | omm | {orientation=row,compression=no}
public | part | table | omm | {orientation=row,compression=no}
public | partsupp | table | omm | {orientation=row,compression=no}
public | region | table | omm | {orientation=row,compression=no}
public | supplier | table | omm | {orientation=row,compression=no}
public | user_dimension | table | omm | {orientation=row,compression=no}
(12 rows)

tpch=#
```

步骤 22 查询 customer 表的数据

```
tpch=# select * from customer limit 10;
```

(截图语句和执行结果)

```
C_custkey | C_name | C_address | C_nationkey | C_phone | C_acctbal | C_kssegment | C_comment
-----+-----+-----+-----+-----+-----+-----+-----
1 | Customer000000001 | Jh2J2qmba et.c.E | 15 | 25-800-741-2988 | 711.56 | BUILDING | to the even, regular platelets, regular, ironic epitaphs hagg
2 | Customer000000002 | XSTf4,McD9ymedstgwfMhchXak | 13 | 23-868-887-8653 | 521.65 | AUTOMOBILE | l accounts, blithely ironic theodolites integrate boldly; caref
3 | Customer000000003 | MGRdIT2w8m | 1 | 11-719-748-1364 | 7498.12 | AUTOMOBILE | deposits eat slyly ironic, even instructions, express foxes detect slyly, blithely even accounts abov
4 | Customer000000004 | xVvS3u4dCn | 4 | 44-128-138-1944 | 2866.93 | HOUSING | requests, final, regular ideas sleep final accou
5 | Customer000000005 | KpvyuPlR84agA(DV6VpZz7f) | 3 | 13-758-942-6364 | 794.47 | HOUSEHOLD | n accounts will have to unwind, foxes cajole accor
6 | Customer000000006 | xZaCmCmPp4dKdRfBv4,LAfRYmWah ym | 28 | 30-114-868-9515 | 7628.57 | AUTOMOBILE | ltem, even deposits book according to the slyly held packages, final accounts cajole requests, furious
7 | Customer000000007 | TCeG5gZgYpP4uSkRvXBfkaodTea | 18 | 28-198-982-9759 | 9561.95 | AUTOMOBILE | ainst the ironic, express theodolites, express, even pinto beans among the exp
8 | Customer000000008 | I8B18d8Ajmc, 8Pv9BPFiyG8ac8Pmml5 | 17 | 27-147-574-9335 | 6819.74 | BUILDING | among the slyly regular theodolites kindly blithely courts, carefully even theodolites haggle slyly along the ide
9 | Customer000000009 | xSL7f0scnfafuMhmfTrj5 | 8 | 18-338-968-8878 | 8324.87 | FURNITURE | r theodolites according to the requests wake thinly excuses, pending requests haggle furtoal
10 | Customer000000010 | 6LEav8KR8PLVcg12Arl Q3rgLzC71 v2 | 5 | 15-741-246-8878 | 2753.54 | HOUSEHOLD | es regular deposits haggle, fur
(10 rows)
```

2. 思考题

数据初始化中出现了 TPC-H，这是什么？

这是一个性能测试程序，主要目的是评价特定查询的决策支持能力，强调服务器在数据挖掘、分析处理方面的能力，用于测试并得到数据库性能指标。

关卡三：openGauss 的 AI4DB 特性应用

1. 关卡验证

(1) 使用 X-Tuner 进行参数优化

步骤 2 在原来 CloudShell 连接窗口中查看 queries01.log。

```
[omm@opengauss01 ~]$ tail -10 /opt/software/tpch-kit/dbgen/queries/queries01.log
```

(截图执行语句和结果)

```
Supplier#000008136          |          3
(411 rows)

 cntrycode | numcust | totacctbal
-----+-----+-----
    13     |    888  | 6737713.99
    17     |    861  | 6460573.72
    18     |    964  | 7236687.40
    23     |    892  | 6701457.95
    29     |    948  | 7158866.63
    30     |    909  | 6808436.13
    31     |    922  | 6806670.18
(7 rows)

total time: 1219390  ms
```

步骤 3 切换至 root 用户，执行 X-Tuner 进行参数建议优化

```
[omm@opengauss01 ~]$ exit
[root@opengauss01 xtuner]# gs_xtuner recommend --db-name tpch --db-user omm --
port 5432 --host 127.0.0.1 --host-user omm
```

(截图执行语句和结果)

name	recommend	min	max	restart
default_statistics_target	1000	100	1000	False
effective_cache_size	21602334	186752	21602334	False
effective_io_concurrency	200	150	250	False
enable_mergejoin	off	0	1	False
enable_nestloop	off	0	1	False
max_connections	370	50	741	True
max_prepared_transactions	370	50	741	True
max_process_memory	28803112	22402420	28803112	True
random_page_cost	1.0	1.0	2.0	False
shared_buffers	186752	186756	214768	True
wal_buffers	5836	2048	5836	True

→ root@opengauss01 /opt/software/openGauss/bin/dbmind/xtuner

步骤 6 获取参数值

```
[omm@opengauss01 ~]$ cd /opt/software/openGauss/data
[omm@opengauss01 data]$ cat postgresql.conf|grep -E 'shared_buffers|
max_connections|effective_cache_size|effective_io_concurrency|wal_buffers|
random_page_cost|default_statistics_target'
```

(截图执行语句和结果)

```
→ omm@opengauss01 /opt/software/openGauss/data cat postgresql.conf|grep -E 'shared_buffers|max_connections|effective_cache
_size|effective_io_concurrency|wal_buffers|random_page_cost|default_statistics_target'
max_connections = 370                                # (change requires restart)
# Note: Increasing max_connections costs ~400 bytes of shared memory per
shared_buffers = 187388                                # min 128kB
bulk_write_ring_size = 2GB                            # for bulkload, max shared_buffers
#standby_shared_buffers_fraction = 0.3                #control shared buffers use in standby, 0.1-1.0
effective_io_concurrency = 200                        # 1-1000; 0 disables prefetching
wal_buffers = 5855                                    # min 32kB
random_page_cost = 1                                # same scale as above
effective_cache_size = 21602940
default_statistics_target = 1000                      # range 1-10000
# max_locks_per_transaction * (max_connections + max_prepared_transactions)
→ omm@opengauss01 /opt/software/openGauss/data
```

步骤 7 再次执行步骤 2，对比优化前的执行时间。

(截图执行语句和结果)

cntrycode	numcust	totacctbal
13	888	6737713.99
17	861	6460573.72
18	964	7236687.40
23	892	6701457.95
29	948	7158866.63
30	909	6808436.13
31	922	6806670.18

(7 rows)

total time: 1182213 ms

步骤 8 【附加题】有兴趣的同学可以尝试并截图记录于此。

(截图执行语句和结果)

(2) Index-advisor: 索引推荐

步骤 4 使用 explain, 对该 SQL 加以分析

```
tpch=# EXPLAIN
SELECT ad.province AS province, SUM(o.actual_price) AS GMV
FROM litemall_orders o,
     address_dimension ad,
     date_dimension dd
WHERE o.address_key = ad.address_key
     AND o.add_date = dd.date_key
     AND dd.year = 2020
     AND dd.month = 3
GROUP BY ad.province
ORDER BY SUM(o.actual_price) DESC;
```

(截图执行语句和结果)

```

QUERY PLAN
-----
Sort (cost=4593.80..4593.88 rows=31 width=47)
  Sort Key: (sum(o.actual_price)) DESC
  -> HashAggregate (cost=4592.72..4593.03 rows=31 width=47)
    Group By Key: ad.province
    -> Hash Join (cost=4354.43..4585.97 rows=1351 width=15)
      Hash Cond: (ad.address_key = o.address_key)
      -> Seq Scan on address_dimension ad (cost=0.00..188.02 rows=8002 width=14)
      -> Hash (cost=4337.54..4337.54 rows=1351 width=9)
        Hash Cond: (o.add_date = dd.date_key)
        -> Seq Scan on litemall_orders o (cost=0.00..3041.00 rows=100000 width=13)
        -> Hash (cost=1031.76..1031.76 rows=2 width=4)
          -> Seq Scan on date_dimension dd (cost=0.00..1031.76 rows=2 width=4)
            Filter: ((year = 2020) AND ((month)::bigint = 3))

(14 rows)

tpch=#

```

步骤 10 使用 explain，对该 SQL 加以分析

```

tpch=# EXPLAIN
SELECT ad.province AS province, SUM(o.actual_price) AS GMV
FROM litemall_orders o,
     address_dimension ad,
     date_dimension dd
WHERE o.address_key = ad.address_key
     AND o.add_date = dd.date_key
     AND dd.year = 2020
     AND dd.month = 3
GROUP BY ad.province
ORDER BY SUM(o.actual_price) DESC;

```

(截图执行语句和结果)

```

QUERY PLAN
-----
Sort (cost=3579.58..3579.65 rows=31 width=47)
  Sort Key: (sum(o.actual_price)) DESC
  -> HashAggregate (cost=3578.50..3578.81 rows=31 width=47)
    Group By Key: ad.province
    -> Hash Join (cost=3340.21..3571.74 rows=1351 width=15)
      Hash Cond: (ad.address_key = o.address_key)
      -> Seq Scan on address_dimension ad (cost=0.00..188.02 rows=8002 width=14)
      -> Hash (cost=3323.32..3323.32 rows=1351 width=9)
        -> Hash Join (cost=17.56..3323.32 rows=1351 width=9)
          Hash Cond: (o.add_date = dd.date_key)
          -> Seq Scan on litemall_orders o (cost=0.00..3041.00 rows=100000 width=13)
          -> Hash (cost=17.53..17.53 rows=2 width=4)
            -> Index Scan using <16503>btree_date_dimension_year on date_dimension dd (cost=0.00..17.53 rows=2 width=4)
              Index Cond: (year = 2020)
              Filter: ((month)::bigint = 3)

(15 rows)

```

步骤 11 【附加题】有兴趣的同学可以尝试并截图记录于此。仅需要从 queries.sql 文件里选择一条或多条进行索引优化即可。

(截图执行语句和结果)

关卡四【附加题】：openGauss 的 DB4AI 特性应用

*本关卡为附加题，有兴趣的同学可以尝试实验并记录于此。

1. 关卡验证

步骤 10 利用训练好的逻辑回归模型预测数据，并与 SVM 算法进行比较，将执行结果截图。

openGauss=# SELECT tax, bath, size, price, price < 100000 AS price_actual, PREDICT BY house_binary_classifier (FEATURES tax, bath, size) AS price_svm_pred, PREDICT BY house_logistic_classifier (FEATURES tax, bath, size) AS price_logistic_pred FROM houses;

(截图执行语句和结果)

```
tpch=# SELECT tax, bath, size, price, price < 100000 AS price_actual, PREDICT BY house_binary_classifier (FEATURES tax, bath, size) AS price_svm_pred, PREDICT BY house_logistic_classifier (FEATURES tax, bath, size) AS price_logistic_pred FROM house
s;
 tax | bath | size | price | price_actual | price_svm_pred | price_logistic_pred
-----+-----+-----+-----+-----+-----+-----
 590 |    1 |   770 | 50000 | t             | t             | t
1050 |    2 |  1410 | 85000 | t             | t             | t
   20 |    1 |  1060 | 22500 | t             | t             | t
   870 |    2 |   1300 | 90000 | t             | t             | t
1320 |    2 |   1500 | 133000 | f             | t             | t
1350 |    1 |    820 | 90500 | t             | f             | f
2790 |   2.5 |  2130 | 260000 | f             | f             | f
   680 |    1 |   1170 | 142500 | f             | t             | t
1840 |    2 |   1500 | 160000 | f             | f             | f
3680 |    2 |  2790 | 240000 | f             | f             | f
1660 |    1 |   1030 | 87000 | t             | f             | f
1620 |    2 |   1250 | 118600 | f             | f             | f
3100 |    2 |   1760 | 140000 | f             | f             | f
2070 |    3 |   1550 | 148000 | f             | f             | f
   650 |   1.5 |   1450 | 65000 | t             | t             | t
(15 rows)

tpch=#
```

清理工作：资源释放

1. 关卡验证

步骤 3 查看到列表中已没有资源时，表示弹性云服务器已删除。

(截图执行语句和结果)

